

Class 11: Structural Bioinformatics pt2

Junlin Ruan (PID: A17839687)

AlphaFold DB

The EBI maintains the largest database of AlphaFold structure prediction models at:
<https://alphafold.ebi.ac.uk>

From last class (before Halloween) we saw that the PDB had 244,290 (Oct 2025)

The total number of protein sequences in UniProtKB is 199,579,901

Key Point: This is a tiny fraction of sequence space that has structural coverage
(0.12%)

```
244290/199579901 * 100
```

```
[1] 0.1224021
```

AFDB is attempting to address this gap...

There are two “Quality Scores” from AlphaFold: one for residues (i.e. each amino acid) called **pLDDT** score, the other **PAE** score measures the confidence in the relative position of the residues (i.e. every pair of residues).

Generating your own structure predictions

Image of all 5 models

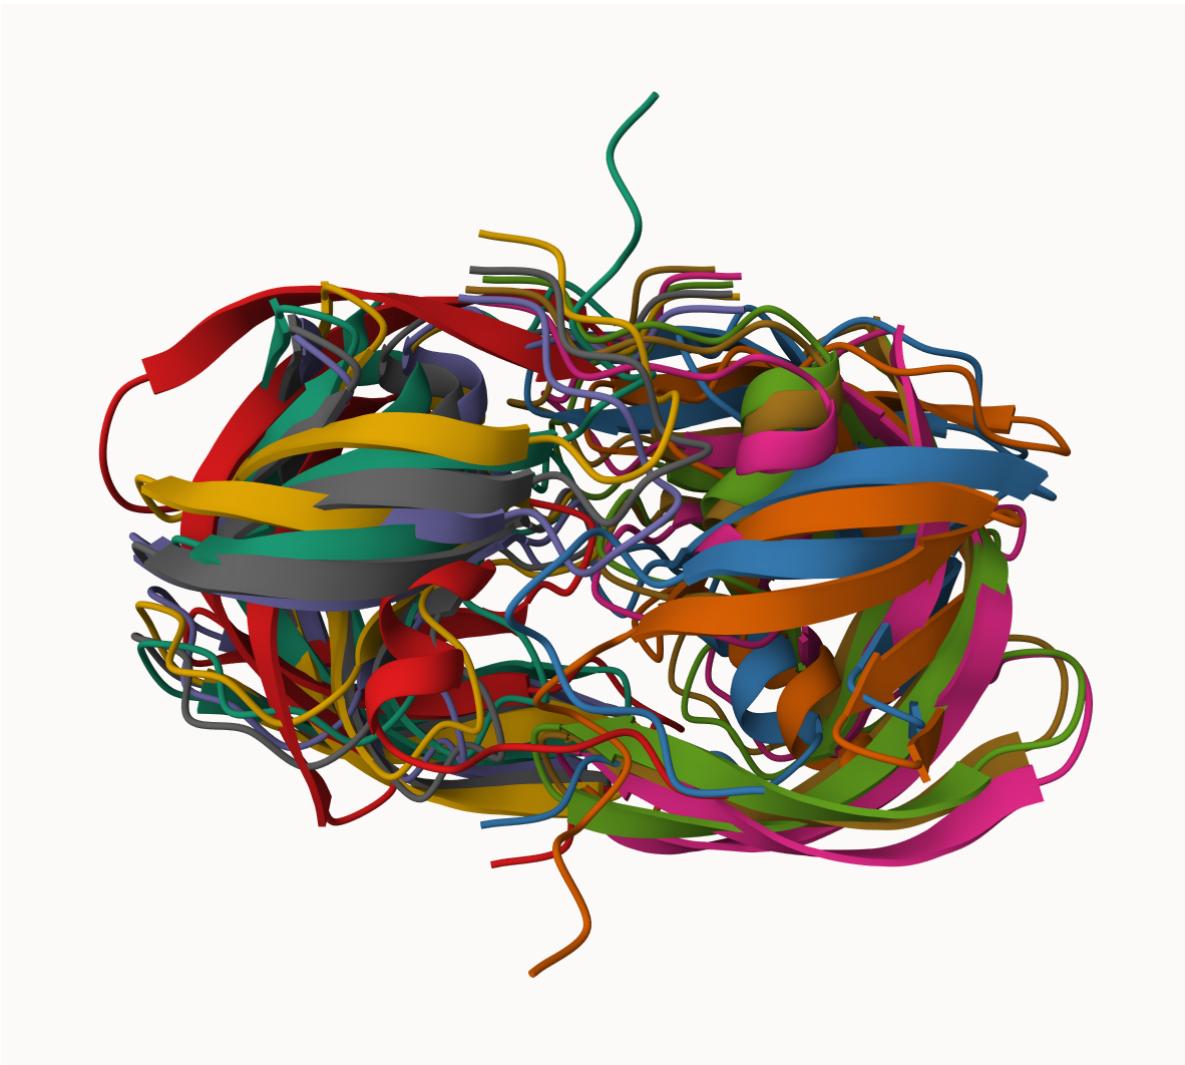
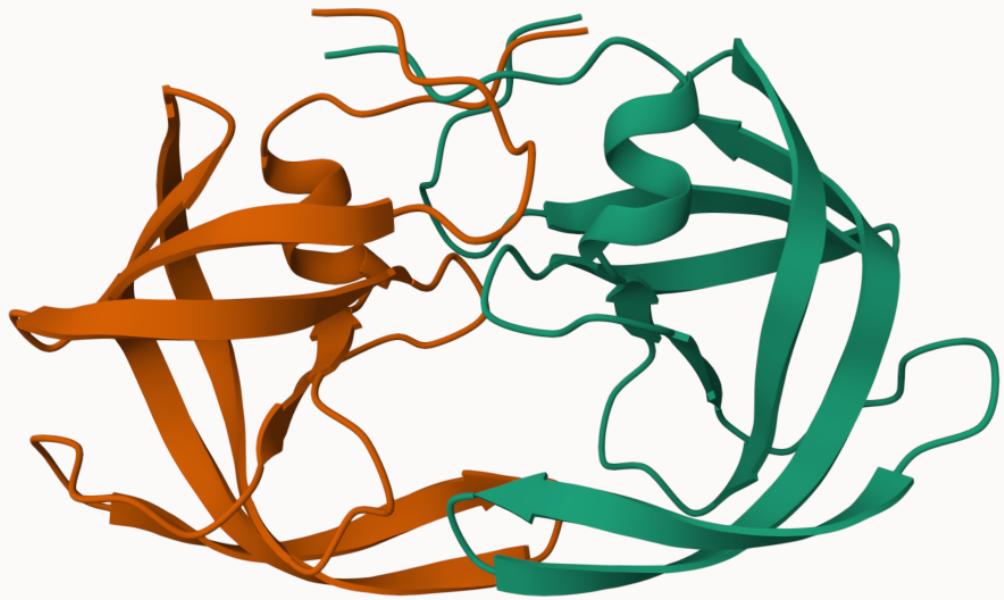
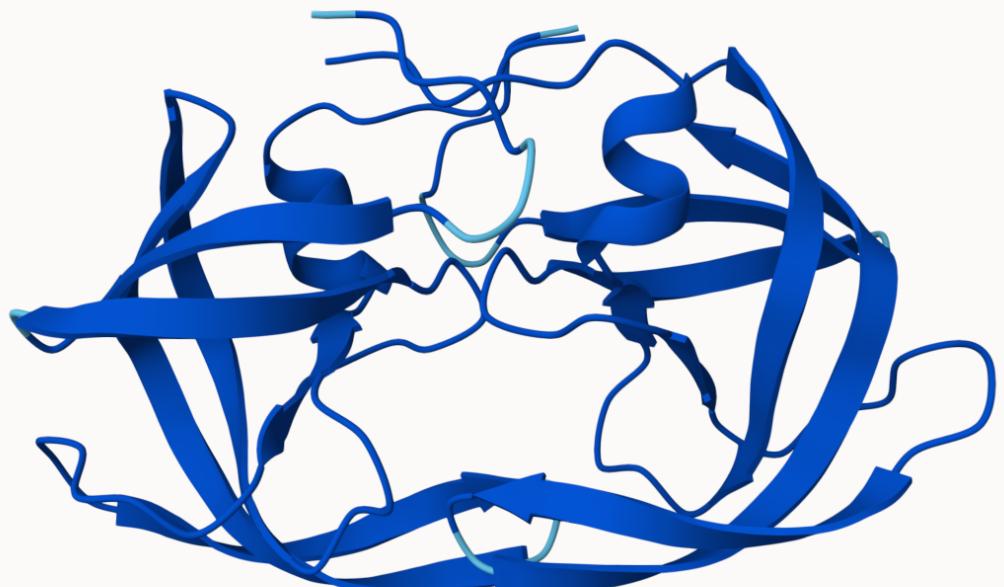


Image of 1st models



pLDDT score of 1st model



pLDDT score of 5th model



Custom analysis of resulting models in R

Read key result files into R. The first thing I need to know is what my results directory/folder is called (i.e. its name is different for every AlphaFold run/job)

```

results_dir <- "HIVPR_Dimer_23119/"

# File names for all PDB models
pdb_files <- list.files(path=results_dir,
                         pattern="*.pdb",
                         full.names = TRUE)

# Print our PDB file names
basename(pdb_files)

```

```

[1] "HIVPR_Dimer_23119_unrelaxed_rank_001_alphaFold2_multimer_v3_model_2_seed_000.pdb"
[2] "HIVPR_Dimer_23119_unrelaxed_rank_002_alphaFold2_multimer_v3_model_4_seed_000.pdb"
[3] "HIVPR_Dimer_23119_unrelaxed_rank_003_alphaFold2_multimer_v3_model_1_seed_000.pdb"
[4] "HIVPR_Dimer_23119_unrelaxed_rank_004_alphaFold2_multimer_v3_model_5_seed_000.pdb"
[5] "HIVPR_Dimer_23119_unrelaxed_rank_005_alphaFold2_multimer_v3_model_3_seed_000.pdb"

```

```

library(bio3d)

m1 <- read.pdb(pdb_files[1])
m1

```

```

Call: read.pdb(file = pdb_files[1])

Total Models#: 1
Total Atoms#: 1514, XYZs#: 4542 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 0 (residues: 0)
Non-protein/nucleic resid values: [ none ]

Protein sequence:
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWPKMIGGIGGFVKVRQYD
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE
ALLDTGADDTVLEEMSLPGRWPKMIGGIGGFVKVRQYDQILIEICGHKAIGTVLVGPTP
VNIIGRNLLTQIGCTLNF

+ attr: atom, xyz, calpha, call

```

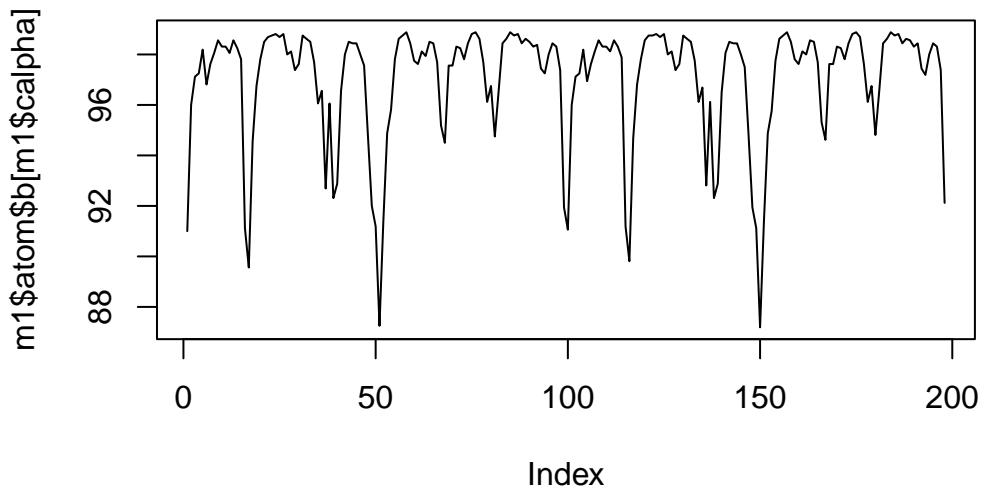
```
head(m1$atom)
```

	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b	segid
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	-17.516	3.072	5.164	1	91	<NA>
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	-17.750	2.633	3.787	1	91	<NA>
3	ATOM	3	C	<NA>	PRO	A	1	<NA>	-16.938	1.399	3.412	1	91	<NA>
4	ATOM	4	CB	<NA>	PRO	A	1	<NA>	-17.328	3.844	2.953	1	91	<NA>
5	ATOM	5	O	<NA>	PRO	A	1	<NA>	-15.969	1.066	4.098	1	91	<NA>
6	ATOM	6	CG	<NA>	PRO	A	1	<NA>	-17.188	4.957	3.941	1	91	<NA>
		elesy	charge											
1		N	<NA>											
2		C	<NA>											
3		C	<NA>											
4		C	<NA>											
5		O	<NA>											
6		C	<NA>											

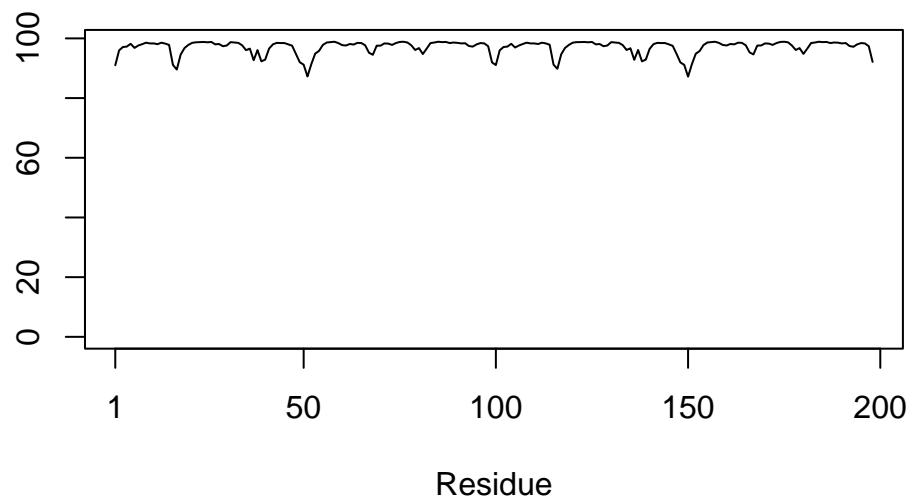
```
m1$atom$b[m1$calpha]
```

```
[1] 91.00 96.00 97.12 97.25 98.19 96.81 97.62 98.06 98.56 98.31 98.31 98.06  
[13] 98.56 98.25 97.81 91.12 89.56 94.56 96.75 97.81 98.50 98.69 98.75 98.81  
[25] 98.69 98.81 98.00 98.12 97.38 97.62 98.75 98.62 98.50 97.69 96.06 96.56  
[37] 92.69 96.06 92.31 92.88 96.56 98.00 98.50 98.44 98.44 98.00 97.56 94.75  
[49] 92.00 91.19 87.25 91.38 94.88 95.81 97.81 98.62 98.75 98.88 98.44 97.75  
[61] 97.62 98.12 97.94 98.50 98.44 97.69 95.19 94.50 97.56 97.56 98.31 98.25  
[73] 97.81 98.44 98.81 98.88 98.62 97.69 96.12 96.75 94.75 96.56 98.44 98.62  
[85] 98.88 98.75 98.81 98.44 98.62 98.50 98.31 98.38 97.44 97.25 98.00 98.44  
[97] 98.31 97.38 91.94 91.06 96.00 97.12 97.25 98.19 96.94 97.62 98.12 98.56  
[109] 98.31 98.31 98.12 98.56 98.31 97.88 91.19 89.81 94.69 96.81 97.81 98.56  
[121] 98.75 98.75 98.81 98.69 98.81 98.00 98.12 97.38 97.62 98.75 98.62 98.50  
[133] 97.75 96.12 96.69 92.81 96.12 92.31 92.88 96.50 98.06 98.50 98.44 98.44  
[145] 98.00 97.50 94.81 91.94 91.12 87.19 91.38 94.88 95.75 97.75 98.62 98.75  
[157] 98.88 98.50 97.81 97.62 98.12 98.00 98.56 98.50 97.69 95.31 94.62 97.62  
[169] 97.62 98.31 98.25 97.81 98.44 98.81 98.88 98.69 97.62 96.12 96.75 94.81  
[181] 96.56 98.44 98.62 98.88 98.75 98.81 98.44 98.62 98.56 98.31 98.44 97.44  
[193] 97.19 98.00 98.44 98.31 97.38 92.12
```

```
plot(m1$atom$b[m1$calpha], typ = "l")
```



```
plot.bio3d(m1$atom$b[m1$calpha], typ = "l")
```



Residue conservation from alignment file

Find the large AlphaFold alignment file

```
aln_file <- list.files(path=results_dir,
                        pattern=".a3m$",
                        full.names = TRUE)
aln_file
```

```
[1] "HIVPR_Dimer_23119//HIVPR_Dimer_23119.a3m"
```

Read this into R

```
aln <- read.fasta(aln_file[1], to.upper = TRUE)
```

```
[1] " ** Duplicated sequence id's: 101 **"
[2] " ** Duplicated sequence id's: 101 **"
```

How many sequences are in this alignment

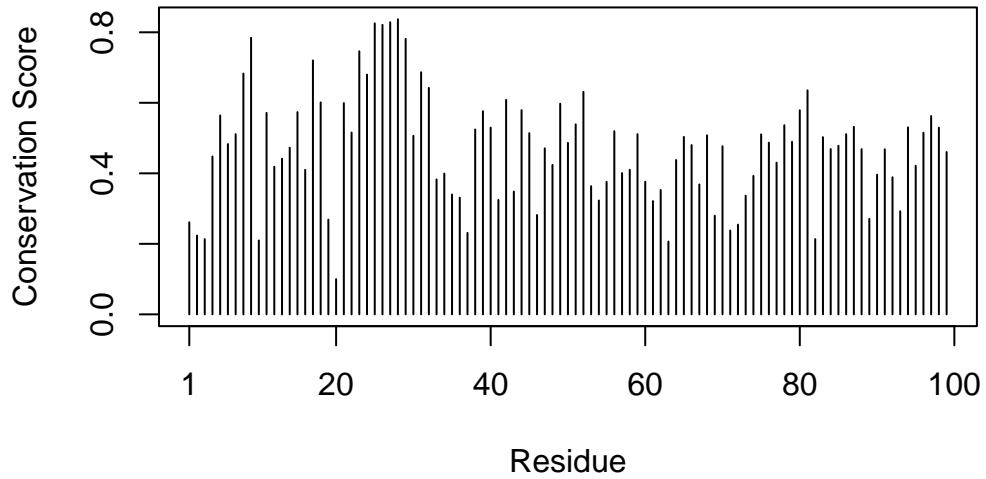
```
dim(aln$ali)
```

```
[1] 5397 132
```

We can score residue conservation in the alignment with the conserv() function.

```
sim <- conserv(aln)
```

```
plotb3(sim[1:99], ylab="Conservation Score")
```



```
con <- consensus(aln, cutoff = 0.9)
con$seq
```

```
[1] "-"
[19] "-"
[37] "-"
[55] "-"
[73] "-"
[91] "-"
[109] "-"
[127] "-"

[1] "D"
[19] "T"
[37] "G"
[55] "A"
[73] "C"
[91] "C"
[109] "C"
[127] "C"
```