

Class12: RNASeq analysis

Junlin Ruan (PID: A17839687)

Table of contents

Background	1
Data Import	1
Toy differential gene expression	3
DESeq2 analysis	8
Volcano Plot	10
Save our results	11
Add gene annotation	11
Pathway analysis	14
Save our main results	16

Background

Today we will analyze some RNASeq data from Himes et al. on the effects of a common steroid (dexamethasone) on airway smooth muscle cells (ASN cells).

The starting point is the “counts” data and “metadata” that contain the count values for each gene in their different experiments (i.e. cell lines with or without the drug).

Data Import

```
# Complete the missing code
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

Let's have a wee peak at these objects:

```
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG000000000003	723	486	904	445	1170
ENSG000000000005	0	0	0	0	0
ENSG000000000419	467	523	616	371	582
ENSG000000000457	347	258	364	237	318
ENSG000000000460	96	81	73	66	118
ENSG000000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		
ENSG000000000003	1097	806	604		
ENSG000000000005	0	0	0		
ENSG000000000419	781	417	509		
ENSG000000000457	447	330	324		
ENSG000000000460	94	102	74		
ENSG000000000938	0	0	0		

Q1. How many genes are in this dataset?

```
nrow(counts)
```

```
[1] 38694
```

Q. How many different experiments (columns in counts or rows in metadata) are there?

```
ncol(counts)
```

```
[1] 8
```

```
nrow(metadata)
```

```
[1] 8
```

```
metadata
```

```

      id      dex celltype      geo_id
1 SRR1039508 control    N61311 GSM1275862
2 SRR1039509 treated    N61311 GSM1275863
3 SRR1039512 control    N052611 GSM1275866
4 SRR1039513 treated    N052611 GSM1275867
5 SRR1039516 control    N080611 GSM1275870
6 SRR1039517 treated    N080611 GSM1275871
7 SRR1039520 control    N061011 GSM1275874
8 SRR1039521 treated    N061011 GSM1275875

```

Q2. How many ‘control’ cell lines we have?

```
sum(metadata$dex == "control")
```

```
[1] 4
```

Toy differential gene expression

To start our analysis let’s calculate the mean counts for all genes in the “control” experiments.

1. Extract all “control” columns from the `counts` object
2. Calculate the mean for all rows (i.e. genes) of these “control” columns
- 3-4. Do the same for “treated” 5. Compare these `control.mean` and `treated.mean` values

```

control.ids <- metadata$dex == "control"
control.counts <- counts[ , control.ids]

```

Q3. How would you make the above code in either approach more robust? Is there a function that could help here?

```
control.means <- rowMeans(control.counts)
```

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called `treated.mean`)

```

treated.ids <- metadata$dex == "treated"
treated.counts <- counts[ , treated.ids]
treated.means <- rowMeans(treated.counts)

```

Store these together for ease of bookkeeping as `meancounts`

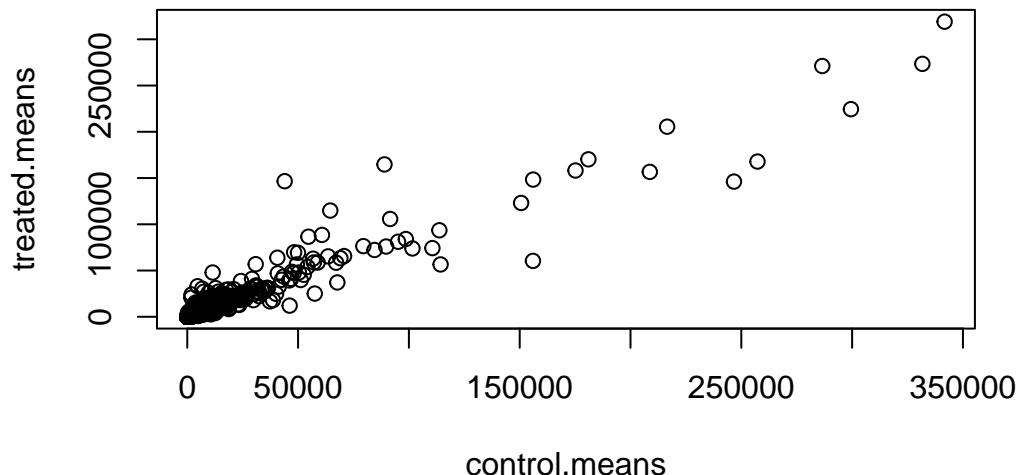
```
meancounts <- data.frame(control.means, treated.means)
head(meancounts)
```

	control.means	treated.means
ENSG000000000003	900.75	658.00
ENSG000000000005	0.00	0.00
ENSG00000000419	520.50	546.00
ENSG00000000457	339.75	316.50
ENSG00000000460	97.25	78.75
ENSG00000000938	0.75	0.00

Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples.

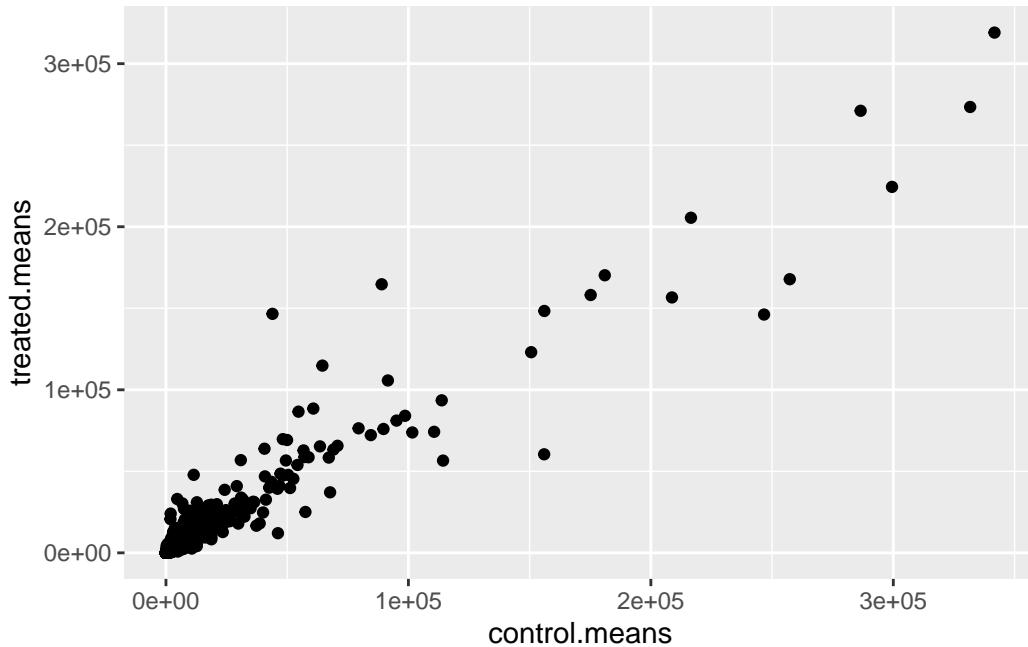
Make a plot of control vs treated mean values for all genes

```
plot(meancounts)
```



Q5 (b). You could also use the ggplot2 package to make this figure producing the plot below. What geom_?() function would you use for this plot?

```
library(ggplot2)
ggplot(meancounts, aes(x = control.means, y = treated.means)) +
  geom_point()
```



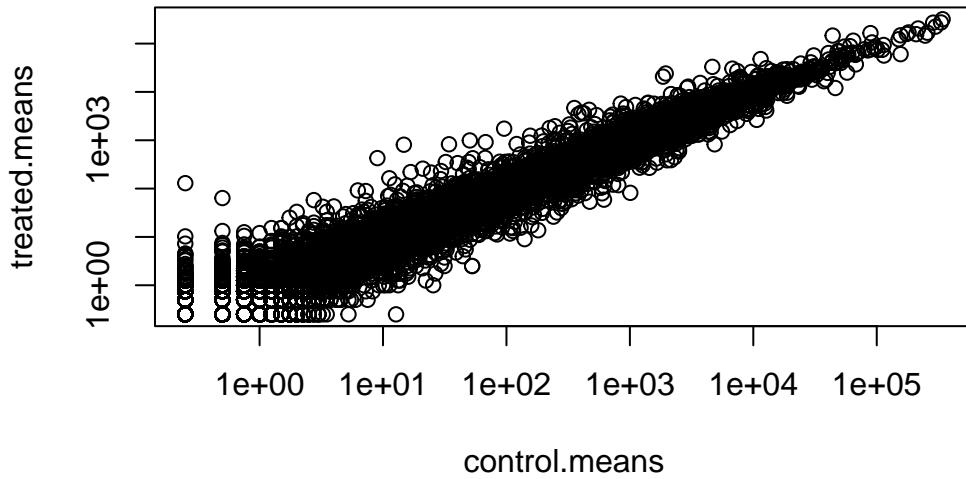
Q6. Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this?

Make this a log log plot

```
plot(meancounts, log = "xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted from logarithmic plot



We often talk metrics like “log2 fold-change”

```
# treated/control
log2(10/10)
```

```
[1] 0
```

```
log2(10/20)
```

```
[1] -1
```

```
log2(20/10)
```

```
[1] 1
```

```
log2(40/10)
```

```
[1] 2
```

```
log2(10/40)
```

```
[1] -2
```

let's calculate the log2 fold change for our treated over control mean counts.

```
meancounts$log2fc <- log2(meancounts$treated.means / meancounts$control.means)
```

```
head(meancounts)
```

	control.means	treated.means	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000938	0.75	0.00	-Inf

```
zero.vals <- which(meancounts[,1:2]==0, arr.ind=TRUE)

to.rm <- unique(zero.vals[,1])
mycounts <- meancounts[-to.rm,]
head(mycounts)
```

	control.means	treated.means	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000971	5219.00	6687.50	0.35769358
ENSG00000001036	2327.00	1785.75	-0.38194109

Q7. What is the purpose of the arr.ind argument in the which() function call above? Why would we then take the first column of the output and need to call the unique() function?

The arr.ind=TRUE argument will cause which() to return both the row and column indices (i.e. positions) where there are TRUE values. Calling unique() will ensure we don't count any row twice if it has zero entries in both samples.

A common “rule of thumb” is a log2 fold change cutoff of +2 and -2 to call genes “Up regulated” or “Down regulated”.

Number of “up” genes

```
sum(meancounts$log2fc >= +2, na.rm = T)
```

```
[1] 1910
```

Number of “down” genes at -2 threshold

```
sum(meancounts$log2fc <= -2, na.rm = T)
```

```
[1] 2330
```

Q8. Using the up.ind vector above can you determine how many up regulated genes we have at the greater than 2 fc level? Q9. Using the down.ind vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

```
up.ind <- meancounts$log2fc > 2  
down.ind <- meancounts$log2fc < (-2)  
sum(up.ind, na.rm = T)
```

```
[1] 1846
```

```
sum(down.ind, na.rm = T)
```

```
[1] 2212
```

Q10. Do you trust these results? Why or why not?

Fold change can be large (e.g. »two-fold up- or down-regulation) without being statistically significant (e.g. based on p-values). So we have to determine if the differences are significant first before drawing conclusions.

DESeq2 analysis

Let’s do this analysis properly and keep our inner stats nerd happy - i.e. are the differences we see between drug and no drug significant given the replicate experiments.

```
library(DESeq2)
```

For DESeq analysis we need three things

- count values (`countData`)
- metadata telling us about the columns in `countData` (`colData`)
- design of the experiment (i.e. what do you want to compare)

Our first function from DESeq2 will setup the input required for analysis by storing all these 3 things together.

```
dds <- DESeqDataSetFromMatrix(countData = counts,
                               colData = metadata,
                               design = ~dex)
```

converting counts to integer mode

```
Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
design formula are characters, converting to factors
```

The main function in DESeq2 that runs the analysis is called `DESeq()`

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```

res <- results(dds)
head(res)

log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 6 columns
  baseMean log2FoldChange    lfcSE     stat   pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG00000000003 747.194195 -0.350703  0.168242 -2.084514 0.0371134
ENSG00000000005  0.000000      NA       NA       NA       NA
ENSG00000000419 520.134160  0.206107  0.101042  2.039828 0.0413675
ENSG00000000457 322.664844  0.024527  0.145134  0.168996 0.8658000
ENSG00000000460 87.682625 -0.147143  0.256995 -0.572550 0.5669497
ENSG00000000938 0.319167 -1.732289  3.493601 -0.495846 0.6200029
  padj
  <numeric>
ENSG00000000003 0.163017
ENSG00000000005  NA
ENSG00000000419 0.175937
ENSG00000000457 0.961682
ENSG00000000460 0.815805
ENSG00000000938  NA

```

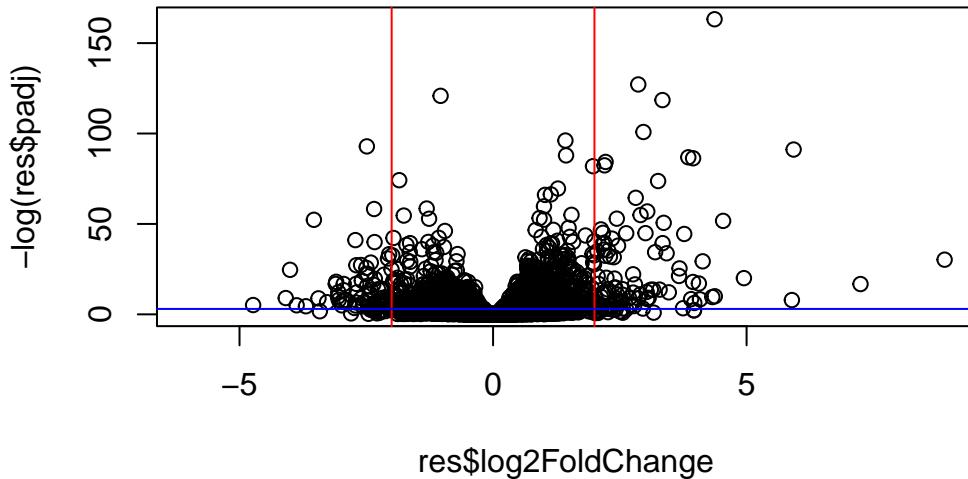
Volcano Plot

This is common summary result figure from these types of experiments and plot the log2 fold-change vs the adjusted p-value.

```

plot(res$log2FoldChange, -log(res$padj))
abline(v = c(-2,2), col = "red")
abline(h = -log(0.05), col = "blue")

```



Save our results

```
write.csv(res, file="my_results.csv")
```

Add gene annotation

To help make sense of our results and communicate them to other folks we need to add some more annotation to our main `res` object.

We will use two bioconductor packages to first map IDs to different formats including the classic gene “symbol” gene name.

I will install these with the following commands: `BiocManager::install("AnnotationDbi")`
`BiocManager::install("org.Hs.eg.db")`

```
library(AnnotationDbi)
library(org.Hs.eg.db)
```

Let's see what is in `org.Hs.eg.db` with the `columns()` function:

```
columns(org.Hs.eg.db)
```

```
[1] "ACCCNUM"      "ALIAS"        "ENSEMBL"       "ENSEMLPROT"    "ENSEMLTRANS"
[6] "ENTREZID"     "ENZYME"       "EVIDENCE"      "EVIDENCEALL"   "GENENAME"
[11] "GENETYPE"     "GO"          "GOALL"         "IPI"          "MAP"
[16] "OMIM"         "ONTOLOGY"     "ONTOLOGYALL"   "PATH"         "PFAM"
[21] "PMID"         "PROSITE"      "REFSEQ"        "SYMBOL"       "UCSCKG"
[26] "UNIPROT"
```

We can translate or “map” IDs between any of these 26 databases using the `mapIds()` function.

```
res$symbol <- mapIds(keys = row.names(res), # our current IDs
                      keytype = "ENSEMBL",    # the format of our IDs
                      x = org.Hs.eg.db,       # where to get the mappings from
                      column = "SYMBOL")      # the format/DB to map to
```

```
'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 7 columns
  baseMean log2FoldChange      lfcSE      stat      pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.350703  0.168242 -2.084514 0.0371134
ENSG000000000005  0.000000      NA        NA        NA        NA
ENSG000000000419 520.134160  0.206107  0.101042  2.039828 0.0413675
ENSG000000000457 322.664844  0.024527  0.145134  0.168996 0.8658000
ENSG000000000460 87.682625 -0.147143  0.256995 -0.572550 0.5669497
ENSG000000000938 0.319167 -1.732289  3.493601 -0.495846 0.6200029
  padj      symbol
  <numeric> <character>
ENSG000000000003 0.163017    TSPAN6
ENSG000000000005  NA          TNMD
ENSG000000000419 0.175937    DPM1
ENSG000000000457 0.961682    SCYL3
ENSG000000000460 0.815805    FIRRM
ENSG000000000938  NA          FGR
```

Add the mappings for “GENENAME” and “ENTREZID” and store as `res$genename` and `res$entrez`

```
res$genename <- mapIds(keys = row.names(res),  
                      keytype = "ENSEMBL",  
                      x = org.Hs.eg.db,  
                      column = "GENENAME")
```

'select()' returned 1:many mapping between keys and columns

```
res$entrez <- mapIds(keys = row.names(res),  
                      keytype = "ENSEMBL",  
                      x = org.Hs.eg.db,  
                      column = "ENTREZID")
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control  
Wald test p-value: dex treated vs control  
DataFrame with 6 rows and 9 columns  
  baseMean log2FoldChange      lfcSE      stat     pvalue  
  <numeric>    <numeric>    <numeric>    <numeric>    <numeric>  
ENSG000000000003 747.194195    -0.350703   0.168242  -2.084514  0.0371134  
ENSG000000000005  0.000000      NA         NA         NA         NA  
ENSG000000000419  520.134160    0.206107   0.101042   2.039828  0.0413675  
ENSG000000000457  322.664844    0.024527   0.145134   0.168996  0.8658000  
ENSG000000000460  87.682625    -0.147143   0.256995  -0.572550  0.5669497  
ENSG000000000938  0.319167    -1.732289   3.493601  -0.495846  0.6200029  
  padj      symbol      genename      entrez  
  <numeric> <character> <character> <character>  
ENSG000000000003  0.163017    TSPAN6       tetraspanin 6      7105  
ENSG000000000005  NA          TNMD        tenomodulin    64102  
ENSG000000000419  0.175937    DPM1 dolichyl-phosphate m..  8813  
ENSG000000000457  0.961682    SCYL3 SCY1 like pseudokina..  57147  
ENSG000000000460  0.815805    FIRRM FIGNL1 interacting r..  55732  
ENSG000000000938  NA          FGR FGR proto-oncogene, ..  2268
```

Pathway analysis

There are lots of bioconductor packages to do this type of analysis. For now let's just try one called **gage** again we need to install this if we don't have it already.

```
library(gage)
library(gageData)
library(pathview)
```

To use **gage** I need two things

- a named vector of fold-change values for our DEGs (our geneset of interest)
- a set of pathways or genesets to use for annotation

```
x <- c("barry" = 5, "lisa" = 10)
x
```

```
barry  lisa
      5     10
```

```
names(x) <- c("low", "high")
x
```

```
low high
 5   10
```

```
foldchanges <- res$log2FoldChange
names(foldchanges) <- res$entrez
head(foldchanges)
```

```
7105       64102       8813       57147       55732       2268
-0.35070296          NA  0.20610728  0.02452701 -0.14714263 -1.73228897
```

```
data("kegg.sets.hs")

keggres = gage(foldchanges, gsets = kegg.sets.hs)
```

In our results object we have:

```

attributes(keggres)

$names
[1] "greater" "less"      "stats"

head(keggres$less, 5)

          p.geomean stat.mean
hsa05332 Graft-versus-host disease 0.0004250607 -3.473335
hsa04940 Type I diabetes mellitus 0.0017820379 -3.002350
hsa05310 Asthma 0.0020046180 -3.009045
hsa04672 Intestinal immune network for IgA production 0.0060434609 -2.560546
hsa05330 Allograft rejection 0.0073679547 -2.501416

          p.val    q.val
hsa05332 Graft-versus-host disease 0.0004250607 0.09053792
hsa04940 Type I diabetes mellitus 0.0017820379 0.14232788
hsa05310 Asthma 0.0020046180 0.14232788
hsa04672 Intestinal immune network for IgA production 0.0060434609 0.31387487
hsa05330 Allograft rejection 0.0073679547 0.31387487

set.size   exp1
hsa05332 Graft-versus-host disease 40 0.0004250607
hsa04940 Type I diabetes mellitus 42 0.0017820379
hsa05310 Asthma 29 0.0020046180
hsa04672 Intestinal immune network for IgA production 47 0.0060434609
hsa05330 Allograft rejection 36 0.0073679547

```

Let's look at one of these pathways (hsa05310 Asthma) with our genes colored up so we can see the overlap

```

pathview(pathway.id = "hsa05310", gene.data = foldchanges)

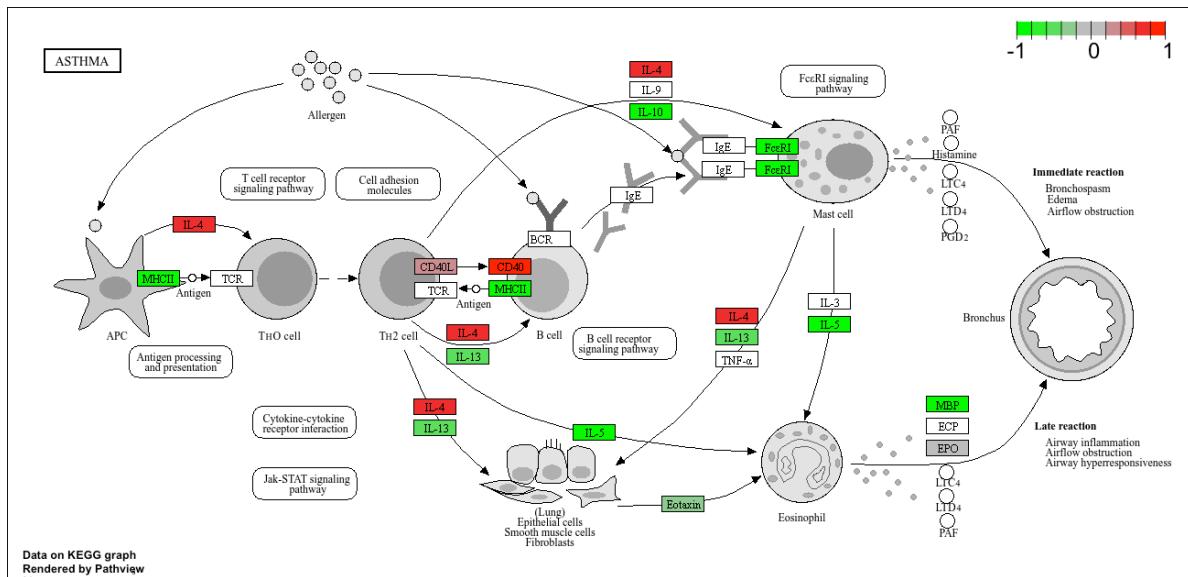
'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/junlinruan/Desktop/courses/BIMM143/class12

Info: Writing image file hsa05310.pathview.png

```

Add this pathway figure to our lab report



Save our main results

```
write.csv(res, file = "myresults_annotated.csv")
```